

Lecture 2

Representational Dimensions

ANNOUNCEMENT

- Assignment 0 due on Wednesday
- You can send it via email by 4:30 on Wednesday. if you want it to count, in case you get into the course

Today's Lecture

 Recap from last lecture

- Representation and Reasoning: Dimensions

- An Overview of This Course
- Further Representational Dimensions
- Intro to search (time permitting)

Course Essentials

- Textbook: Artificial Intelligence: Foundations of Computational Agents. by Poole and Mackworth. (P&M)
- We will cover at least Chapters: 1, 3, 4, 5, 6, 8, 9
- Connect for assignments and marks
- Piazza for Discussion board
- Alspace: online tools for learning Artificial Intelligence
<http://aispace.org/>

What is Artificial Intelligence?

Question: We use the following definition

- The study and design of
 - A. Systems that think rationally
 - B. Systems that act like humans
 - C. Systems that act rationally
 - D. Systems that think like humans

What is Artificial Intelligence?

Clicker Question: We use the following definition

- The study and design of
 - A. Systems that think rationally
 - B. Systems that act like humans
 - C. Systems that act rationally
 - D. Systems that think like humans

AI as Study and Design of Intelligent Agents

- **Intelligent agents**: artifacts that **act rationally** in their environment
- Their **actions** are **appropriate** for their goals and circumstances
- They are **flexible** to changing environments and goals
- They **learn** from experience
- They make **appropriate choices** given **perceptual limitations** and **limited resources**
- This definition drops the constraint of cognitive plausibility

- Same as building flying machines by understanding general principles of flying (aerodynamic) vs. by reproducing how birds fly

Knowledge Representation
Machine Learning

abilities

prior knowledge

past experiences

goals/values

observations

Agent

Representation
& Reasoning

Reasoning +
Decision Theory

Actions

Natural Language
Generation

+
Robotics

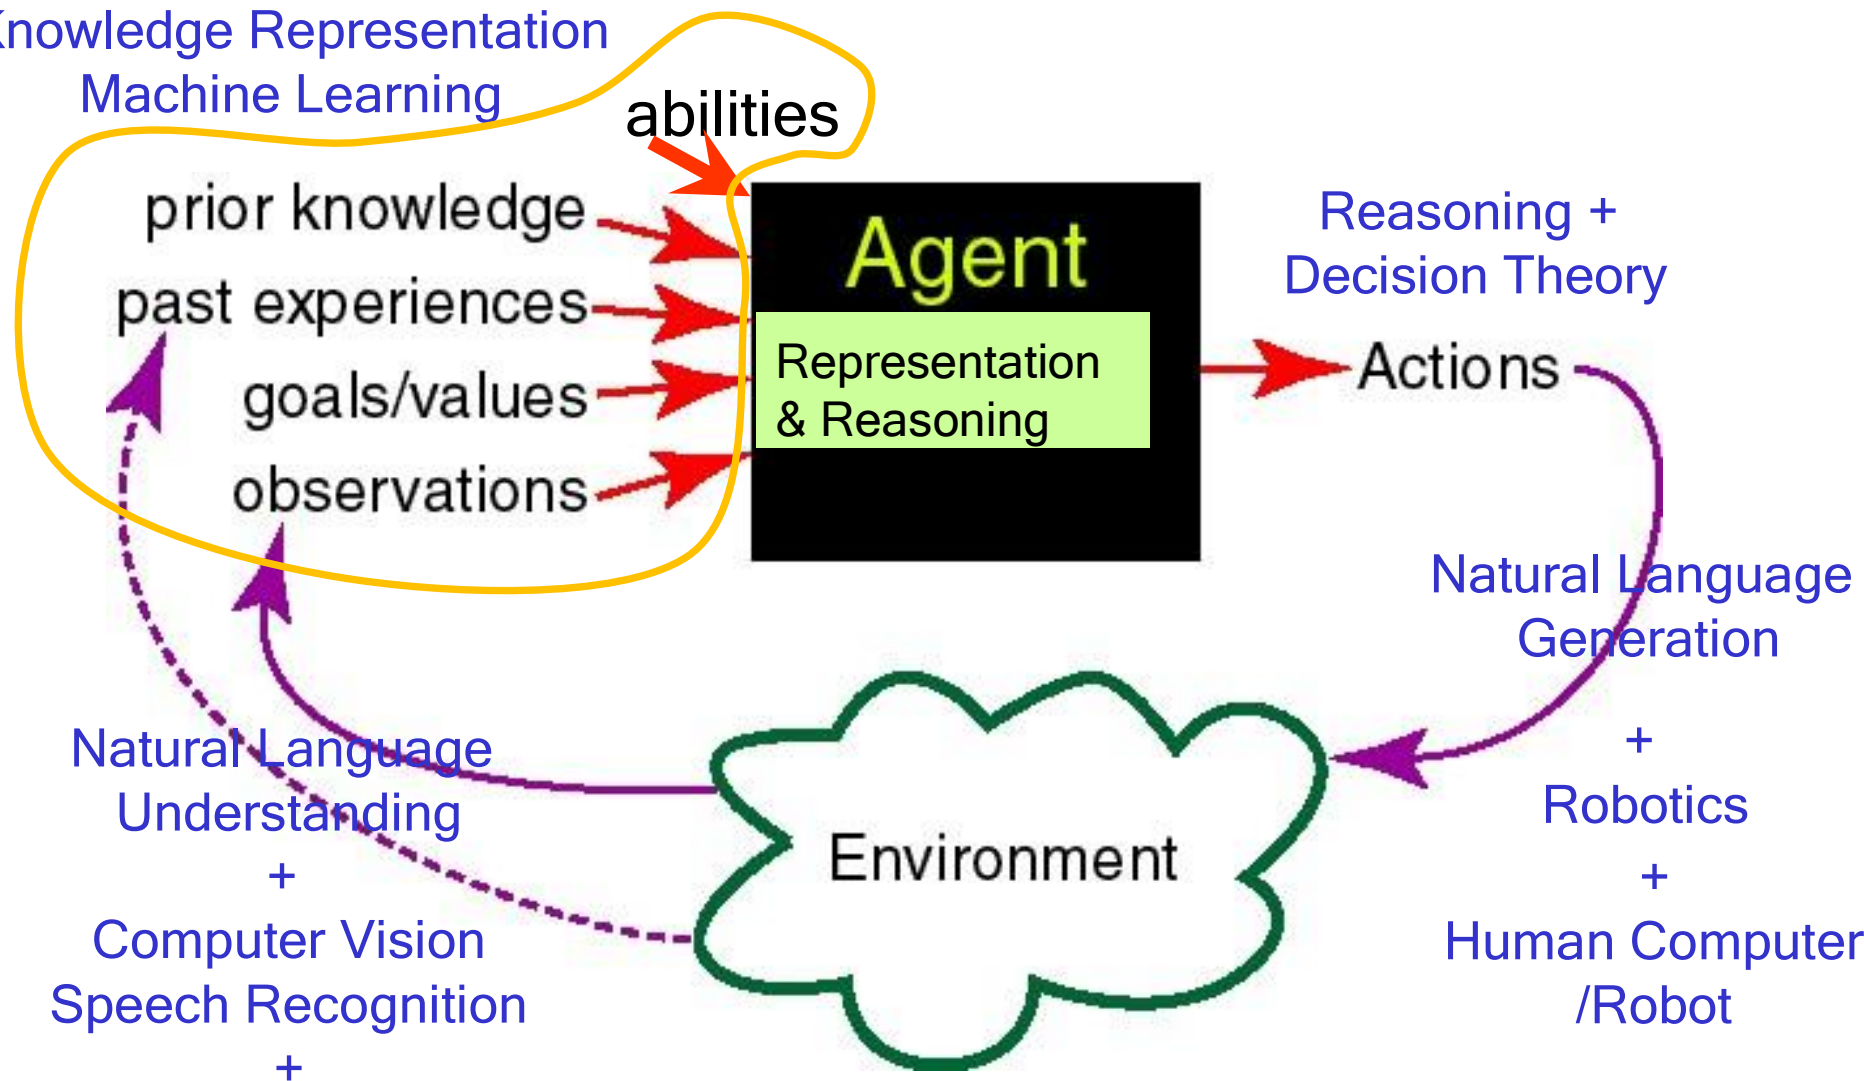
+
Human Computer
/Robot

Environment

Natural Language
Understanding

+
Computer Vision
Speech Recognition


+
Intelligent Agents in the World



Today's Lecture

- Recap from last lecture
- ➡ • Representation and Reasoning: Dimensions
- An Overview of This Course
- Further Representational Dimensions

- Intro to search (time permitting)



Representation and Reasoning

Representation
& Reasoning

To use these inputs an agent needs to represent them

□ knowledge

One of AI goals: specify how a system can

- Acquire and represent knowledge about a domain
(representation)
- Use the knowledge to solve problems in that domain
(reasoning)

Representation and Reasoning (R&R) System

- A **representation language** to describe
- The environment

- Problems (questions/tasks) to be solved
- Computational **reasoning procedures** to compute a solution to a problem
- E.g., an answer, sequence of actions problem \Rightarrow representation

\Rightarrow computation \Rightarrow representation \Rightarrow solution

- Choice of an **appropriate R&R system** depends on various dimensions, e.g. properties of
- the environment, the type of problems, the agent, the computational resources, etc.

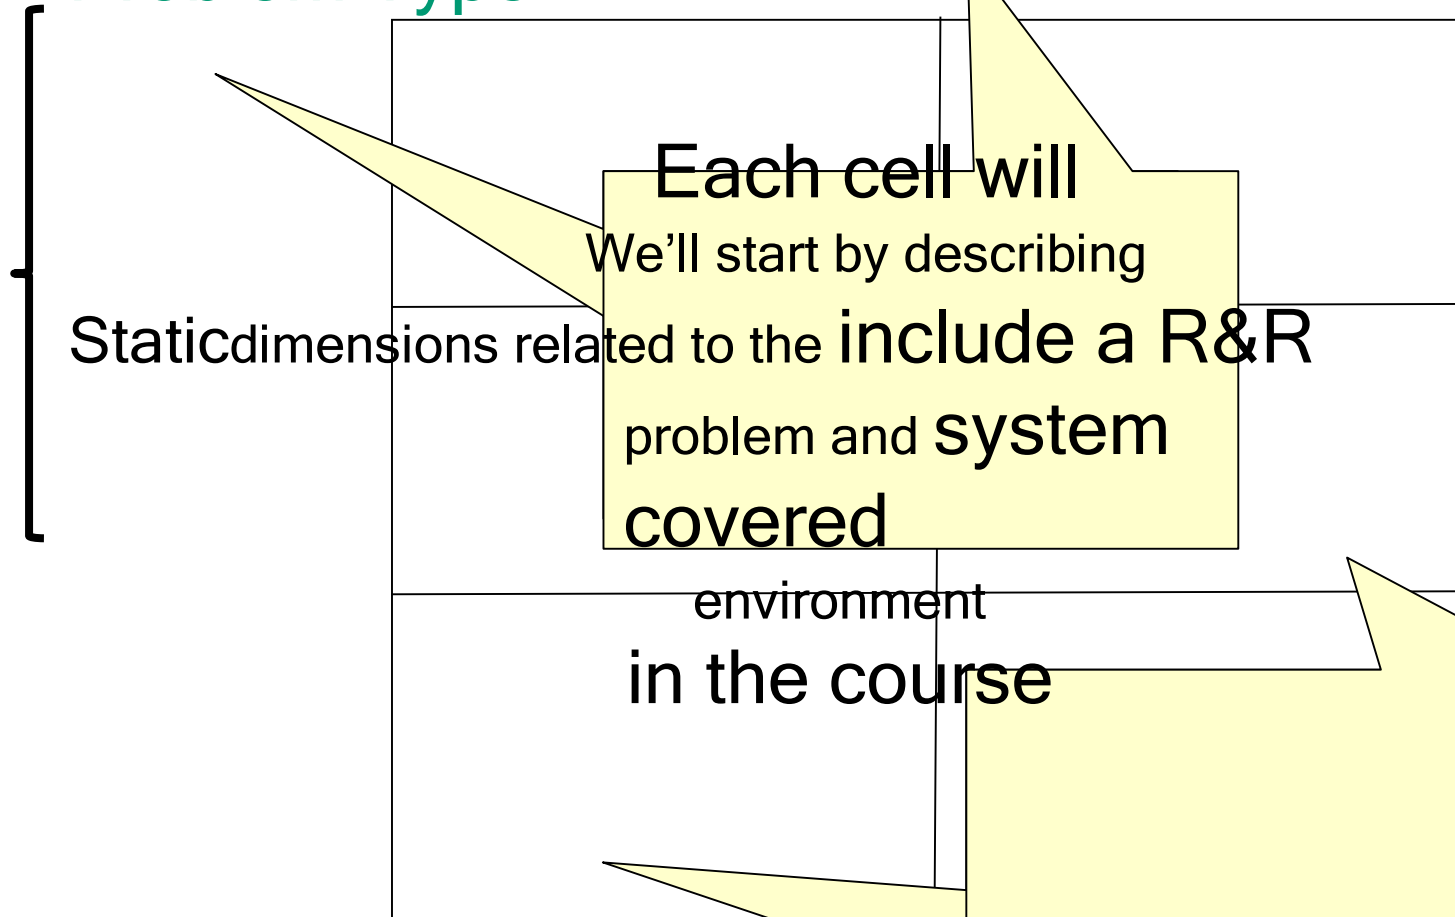
Representational Dimensions

Environment

Deterministic

Stochastic

Problem Type



Sequential

Then we'll include in each
cell the various R&R Then
we'll include in systems
covered in the each cell
R&R system

course, and discuss some covered in the course
more dimensions

Problem Types

- **Static:** finding a solution does not involve reasoning into the future (time is ignored)
- One-step solution

- **Sequential:** finding a solution requires looking for a number of steps into the future, e.g.,
- Fixed horizon (fixed number of steps)
- Indefinite horizon (finite, but unknown number of steps)

Problem Types

- **Constraint Satisfaction** - Find state that satisfies set of constraints (**static**).
- e.g., what is a feasible schedule for final exams?
- **Answering Query** - Is a given proposition true/likely given what is known? (**static**).
- e.g., does the patient suffers from viral hepatitis?

- **Planning** - Find sequence of actions to reach a goal state / maximize outcome (sequential).
- e.g., Navigate through an environment to reach a particular location

Deterministic vs. Stochastic (Uncertain) Environment Representational Dimensions

Environment

Deterministic

Stochastic

Problem Type

Static
Constraint
Satisfaction

Query

Sequential

Planning

Deterministic vs. Stochastic (Uncertain) Environment

- **Sensing Uncertainty:** The agent cannot fully observe the current state of the world when acting

	Sensing Uncertainty?
Teacher's explanation	
Soccer Player Kick	

- **Effect Uncertainty:** the agent does not know for sure the immediate effects of its actions

Deterministic vs. Stochastic (Uncertain) Environment

- **Sensing Uncertainty:** The agent cannot fully observe the current state of the world

	Sensing Uncertainty?
Teacher's explanation	YES
Soccer Player Kick	NO

- **Effect Uncertainty:** the agent does not know for sure the effects of its actions

Deterministic vs. Stochastic (Uncertain) Environment

	Effect Uncertainty?
Teacher's explanation	
Soccer Player Kick	

- **Sensing Uncertainty:** The agent cannot fully observe the current state of the world

	Sensing Uncertainty?
Teacher's explanation	YES
Soccer Player Kick	NO

- **Effect Uncertainty**: the agent does not know for sure the effects of its actions

	Effect Uncertainty?
Teacher's explanation	YES
Soccer Player Kick	YES

Question: Chess and Poker

An environment is **stochastic** if at least one of these is true

- Sensing Uncertainty: the agent cannot fully observe the current state of the world

Deterministic vs. Stochastic (Uncertain) Environment

- Effect Uncertainty: the agent does not know for sure the **immediate, direct** effects of its actions
- A. Poker and Chess are both stochastic
- B. Chess is stochastic and Poker is deterministic
- C. Poker and Chess are both deterministic
- D. Chess is deterministic and Poker is stochastic

Question: Chess and Poker

An environment is **stochastic** if at least one of these is true

- Sensing Uncertainty: the agent cannot fully observe the current state of the world
- Effect Uncertainty: the agent does not know for sure the **immediate, direct** effects of its actions

A. Poker and Chess are both stochastic

B. Chess is stochastic and Poker is deterministic

C. Poker and Chess are both stochastic

D. Chess is deterministic and Poker is stochastic

Deterministic vs. Stochastic Domains

- Historically, AI has been divided into two camps: those who prefer representations based on **logic** and those who prefer **probability**.

Note: Some of the most exciting current research in AI is actually building bridges between these camps.

Representational Dimensions

Environment

Deterministic

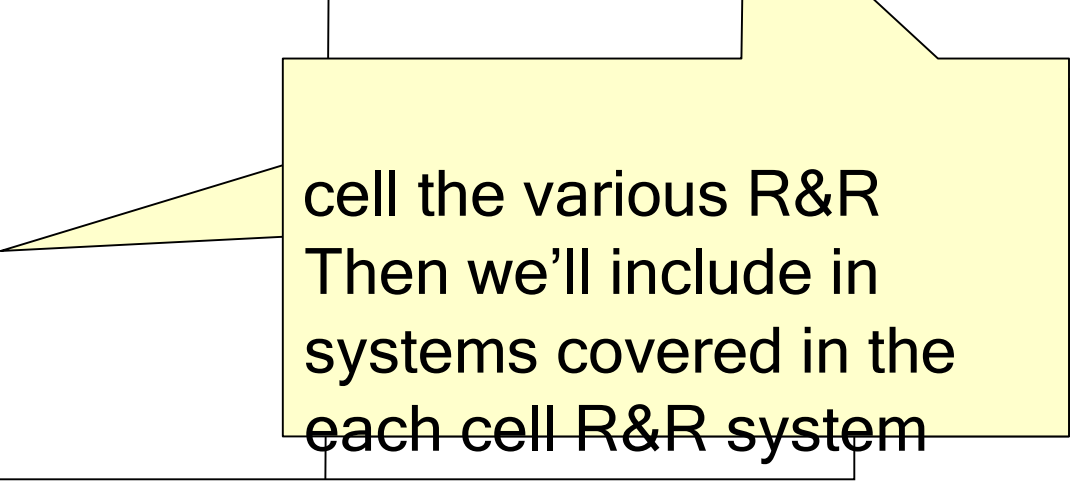
Stochastic

Problem Type

We described Each cell will
Static dimensions include a R&R
related to the system covered
problem and in the course environment

Sequential

Now we include in each




cell the various R&R
Then we'll include in
systems covered in the
each cell R&R system

course, and discuss some covered in the course
more dimensions

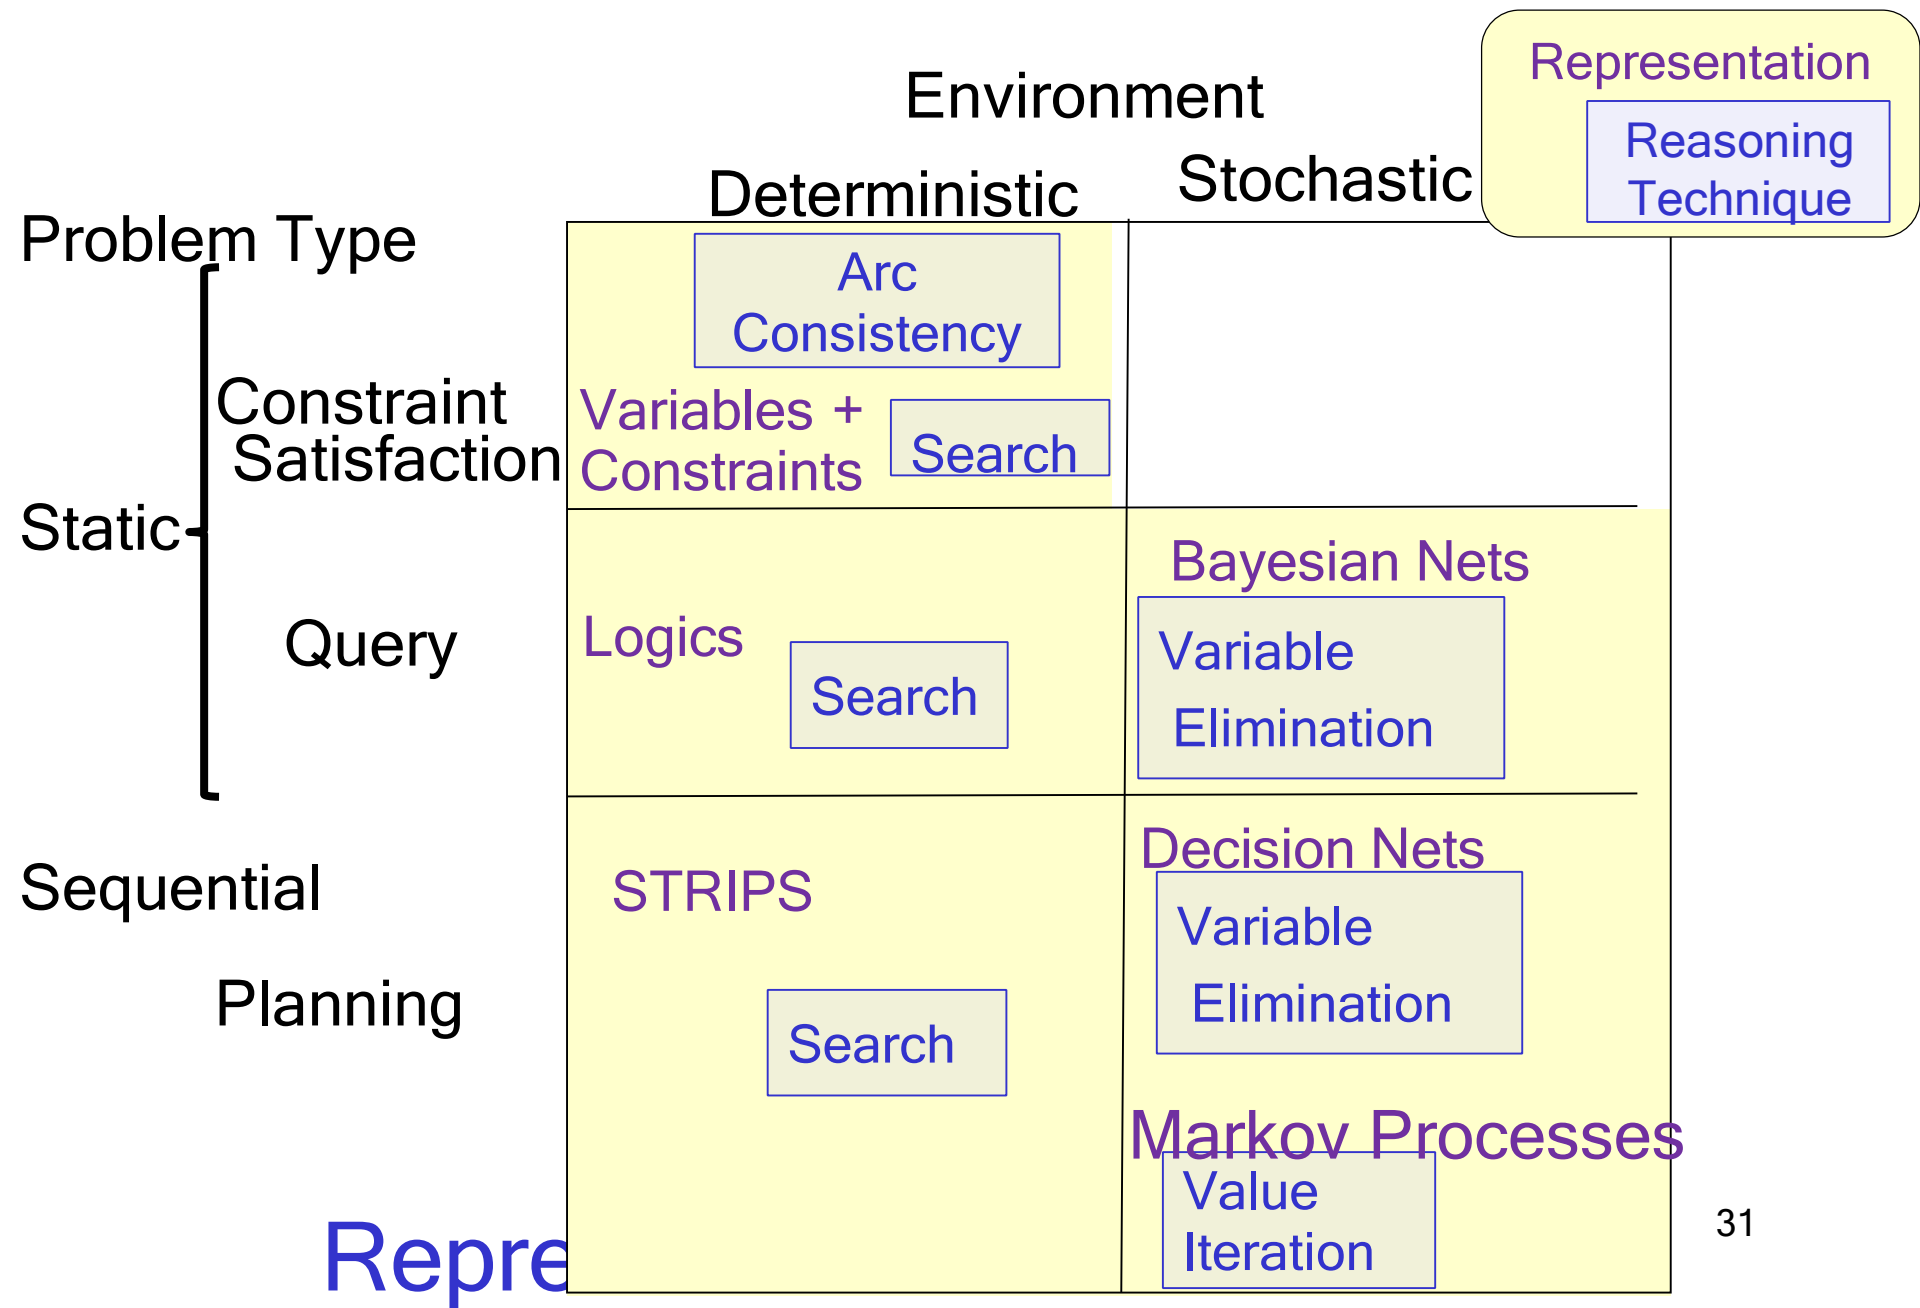
Today's Lecture

- Recap from last lecture
- Representation and Reasoning: Dimensions



An Overview of This Course

- Further Representational Dimensions
- Intro to search (time permitting)



Other Representational Dimensions

We've already discussed:

- Problem Types (**Static** vs. **Sequential**)
- **Deterministic** versus **stochastic** domains

Some other important dimensions

- Representation scheme: **Explicit state** or **features** or **relations**
- **Flat** or **hierarchical** representation
- **Knowledge given** versus **knowledge learned** from experience
- **Goals** versus **complex preferences**

- Single-agent vs. multi-agent

Today's Lecture

- Recap from last lecture
- Representation and Reasoning: Dimensions
- An Overview of This Course
- ➡ • Further Representational Dimensions
- Intro to search (time permitting)

Explicit State vs Features

How do we model the environment?

- You can enumerate the **states** of the world or
- A state can be described in terms of **features**
- Often a more natural description
- 30 binary features (also called propositions) can represent



How do we model the environment?

Explicit State vs Features

- You can enumerate the **states** of the world or
- A state can be described in terms of **features**
- Often a more natural description
- 30 binary features (also called propositions) can represent

$$2^{30}=1,073,741,824 \text{ states}$$

Mars Explorer

Weather

{S, C}

Example

Explicit State vs Features

Temperature [-40, 40]

Longitude

[0, 359]

[0, 179]

Latitude

One possible state

Number of possible states (mutually exclusive)

Mars Explorer Example

Explicit State vs Features

Weather

$\{S, C\}$

Temperature

$[-40, 40]$

Longitude

$[0, 359]$

$[0, 179]$

Latitude

One possible state

$\{S, -30, 320, 210\}$

Number of possible states (mutually exclusive)

$2 \times 81 \times 360 \times 180$

Explicit State vs. Features vs. Relations

- States can be described in terms of **objects** and **relationships**
- There is a proposition for each relationship on each tuple of objects
- University Example:
 - Students (S) = {s1, s2, s3, ..., s200}
 - Courses (C) = {c1, c2, c3, ..., c10}
 - Registered (S, C)
- Number of Relations:

- Number of Propositions:
- Number of States:

Explicit State vs. Features vs. Relations

- States can be described in terms of **objects** and **relationships**
- There is a proposition for each relationship on each tuple of objects
- University Example:
- Students (S) = {s1, s2, s3, ..., s200}

- Courses (C) = {c1, c2, c3, ..., c10}
- Registered (S, C)
- Number of Relations: 1 • Number of Propositions:

$$200 * 10$$

33

- Number of States:

$$2^{2000}$$

Question

One binary relation (e.g., likes) and 9 individuals (e.g. people). How many states?

40

A. 81^2

B. 2^{18}

C. 2^{81}

D. 10^9

Question

One binary relation (e.g., likes) and 9 individuals (e.g. people). How many states?

A. 81^2

B. 2^{18}

C. 2^{81}

D. 10^9

Flat vs. hierarchical

- Should we model the whole world on the same level of abstraction?
- Single level of abstraction: **flat**
- Multiple levels of abstraction: **hierarchical**
- Example: Planning a trip from here to a resort in
Cancun

Going to the airport

Take a cab

Call a cab

Lookup number

Dial number

Ride in the cab

Check in

Pay for the cab

- This course: mainly **flat representations**
- Hierarchical representations required for scaling up.

Knowledge given vs. knowledge learned from experience

- The agent is provided with a model of the world once and for all **or**
- The agent **can learn** how the world works based on experience

- in this case, the agent often still needs some prior knowledge
- This course: mostly **knowledge given**
- **Learning:** CPSC 340 and CPSC 422

Goals vs. (complex) preferences

- An agent may have a **goal** that it wants to achieve, e.g.,
- there is some **state** or **set of states** that the agent wants to be in
- there is **some proposition** or **set of propositions** that the agent wants to make true
- An agent may have **preferences**

- a **preference/utility function** describes how happy the agent is in each state of the world
- Agent's task is to reach a state which makes it as happy as possible
- Preferences can be **complex**
 - What beverage to order?
- I am in a hurry so I need something quickly
- I like Cappuccino better than Espresso, but it takes longer to make...
- This course: **goals** and **simple preferences**

Single-agent vs. Multi-agent domains

- Does the environment include other agents?

- If there are other agents, it can be useful to explicitly model
 - their goals and beliefs,
 - how they react to our actions
- Other agents can be: cooperative, competitive, or a bit of both
- This course: only **single agent scenario**

Summary

Would like most general agents possible, but in this course we need to restrict ourselves to:

- **Flat** representations (vs. hierarchical)
- **Knowledge given** (vs. knowledge learned)

- Goals and simple preferences (vs. complex preferences)
- Single-agent scenarios (vs. multi-agent scenarios)

We will look at


- Deterministic and stochastic domains
- Static and Sequential problems

And see examples of representations using

- Explicit state or features or relations

AI Application

- At the beginning of next class, we will look at some AI applications that you have found for your assignment 0



- You are asked to described them in terms of the elements above and some more
- What does the AI application do
- Goals

- prior knowledge needed
- past experiences that it does (or could) learn from
- Observations needed
- Actions performed
- AI technologies used
- Why is it intelligent?
- Evaluation?

Today's Lecture

- Recap from last lecture
- Representation and Reasoning: Dimensions
- An Overview of This Course
- Further Representational Dimensions
- ➡ • Intro to Search (time permitting)

Representational Dimensions

Environment

Representation

Reasoning
Technique

Deterministic

Stochastic

Problem Type

Constraint
Satisfaction

Static

Query

Sequential

Planning

Arc
Consistency

Variables +
Constraints

Search

Logics

Search

Bayesian Nets

Variable
Elimination

STRIPS

Search

Decision Nets

Variable
Elimination

Representational Dimensions

Environment

Representation

Reasoning
Technique

Deterministic

Stochastic

Problem Type

Constraint
Satisfaction

Static

Query

Sequential

Planning

First Part of
the Course

Arc
Consistency

Variables +
Constraints

Search

Logics

Search

Bayesian Nets

Variable
Elimination

STRIPS

Search

Decision Nets

Variable
Elimination

Representational Dimensions

Environment

Representation

Reasoning
Technique

Deterministic

Stochastic

Problem Type

Constraint
Satisfaction

Static

Query

Sequential

Planning

We'll focus
on Search

Arc
Consistency

Variables +
Constraints

Search

Logics

Search

Bayesian Nets

Variable
Elimination

STRIPS


Search

Decision Nets

Variable
Elimination

(Adversarial) Search: Checkers

- Early learning work in 1950s by Arthur Samuel at IBM
- Chinook program by Jonathan Schaeffer (UofA)
- Search explores the space of possible moves and their consequence



Source: *IBM Research*

- ✓ 1994: world champion
- ✓ 2007: declared unbeatable

(Adversarial) Search: Chess

In 1997, **Gary Kasparov**, the chess grandmaster and reigning world champion played against **Deep Blue**, a program written by researchers at IBM



Source: *IBM Research*



Slide

(Adversarial) Search: Chess

- Deep Blue's won 3 games, lost 2, tied 1

May 11th, 1997

Computer won world champion of chess

(Deep Blue)

(Garry Kasparov)




- 30 CPUs + 480 chess processors
- Searched 126 million states per sec
- Generated 30 billion positions per move reaching depth 14 routinely

(Reuters = Kyodo News)

Slide

Search

- Often we are not given an algorithm to solve a problem, but only a specification of what a solution is 

we have to **search** for a solution.

- Enumerate a set of potential partial solutions
- Check to see if they are solutions or could lead to one

Simple Search Agent

Deterministic, goal-driven agent

- Agent is in a **start state**
- Agent is given a **goal** (subset of possible states)
- Environment changes only when the agent acts
- Agent perfectly knows:
 - **actions** that can be applied in any given state
 - the **state** it is going to end up in when an action is applied in a given state

- The sequence of actions (and appropriate ordering) taking the agent from the start state to a goal state is the **solution**

Definition of a search problem

- **Initial state(s)**
- Set of **actions** (operators) available to the agent
- An **action function** that, given a state and an action, returns a new state
- **Goal state(s)**
- **Search space**: set of states that will be searched for a path from initial state to goal, given the available actions
- **states are nodes** and **actions are links** between them.

- Not necessarily given explicitly (state space might be infinite)
- **Path Cost** (we ignore this for now)

Three examples

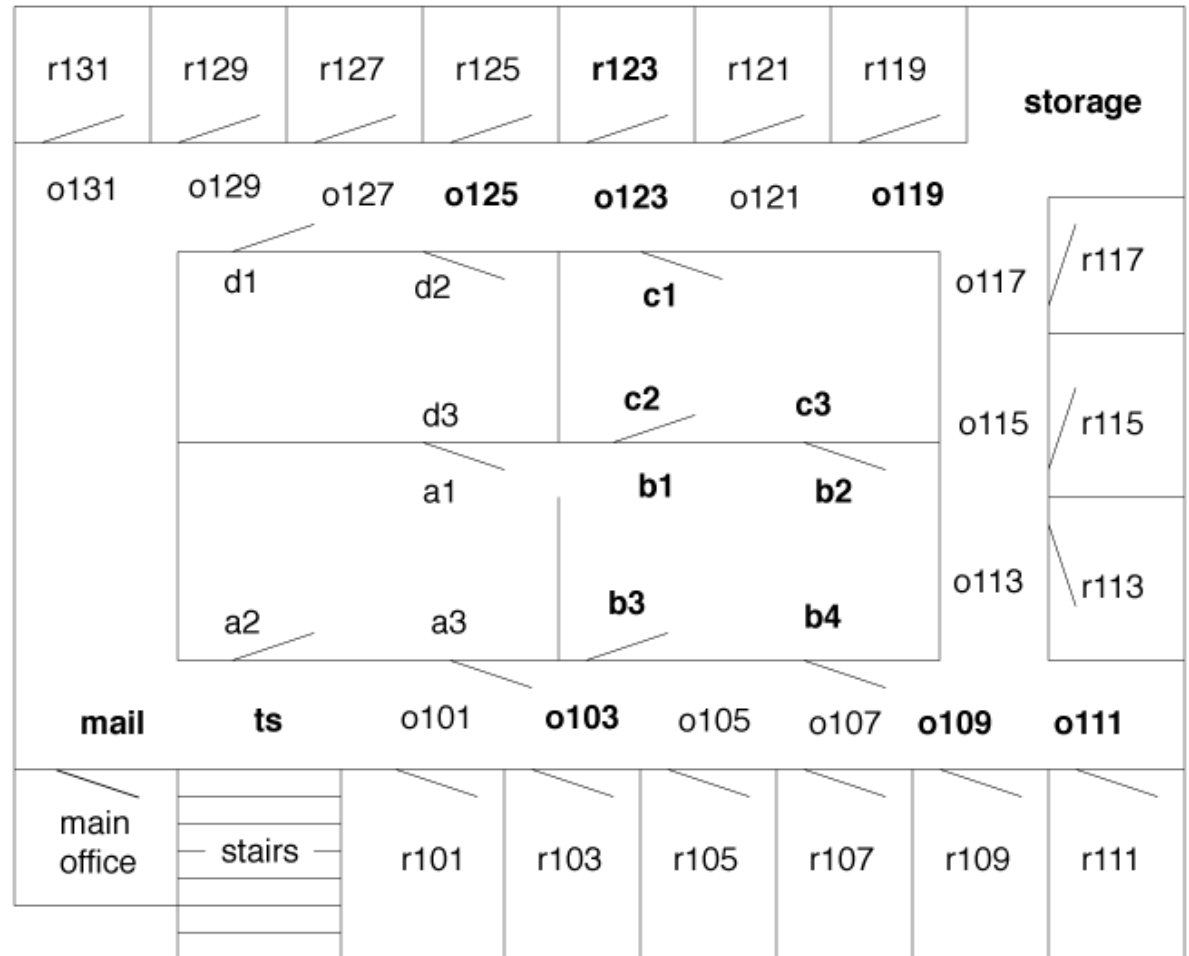
1. The delivery robot planning the route it will take in a building to get from one room to another (Ch 1.6)
2. Vacuum cleaner world
3. Solving an 8-puzzle

Slide

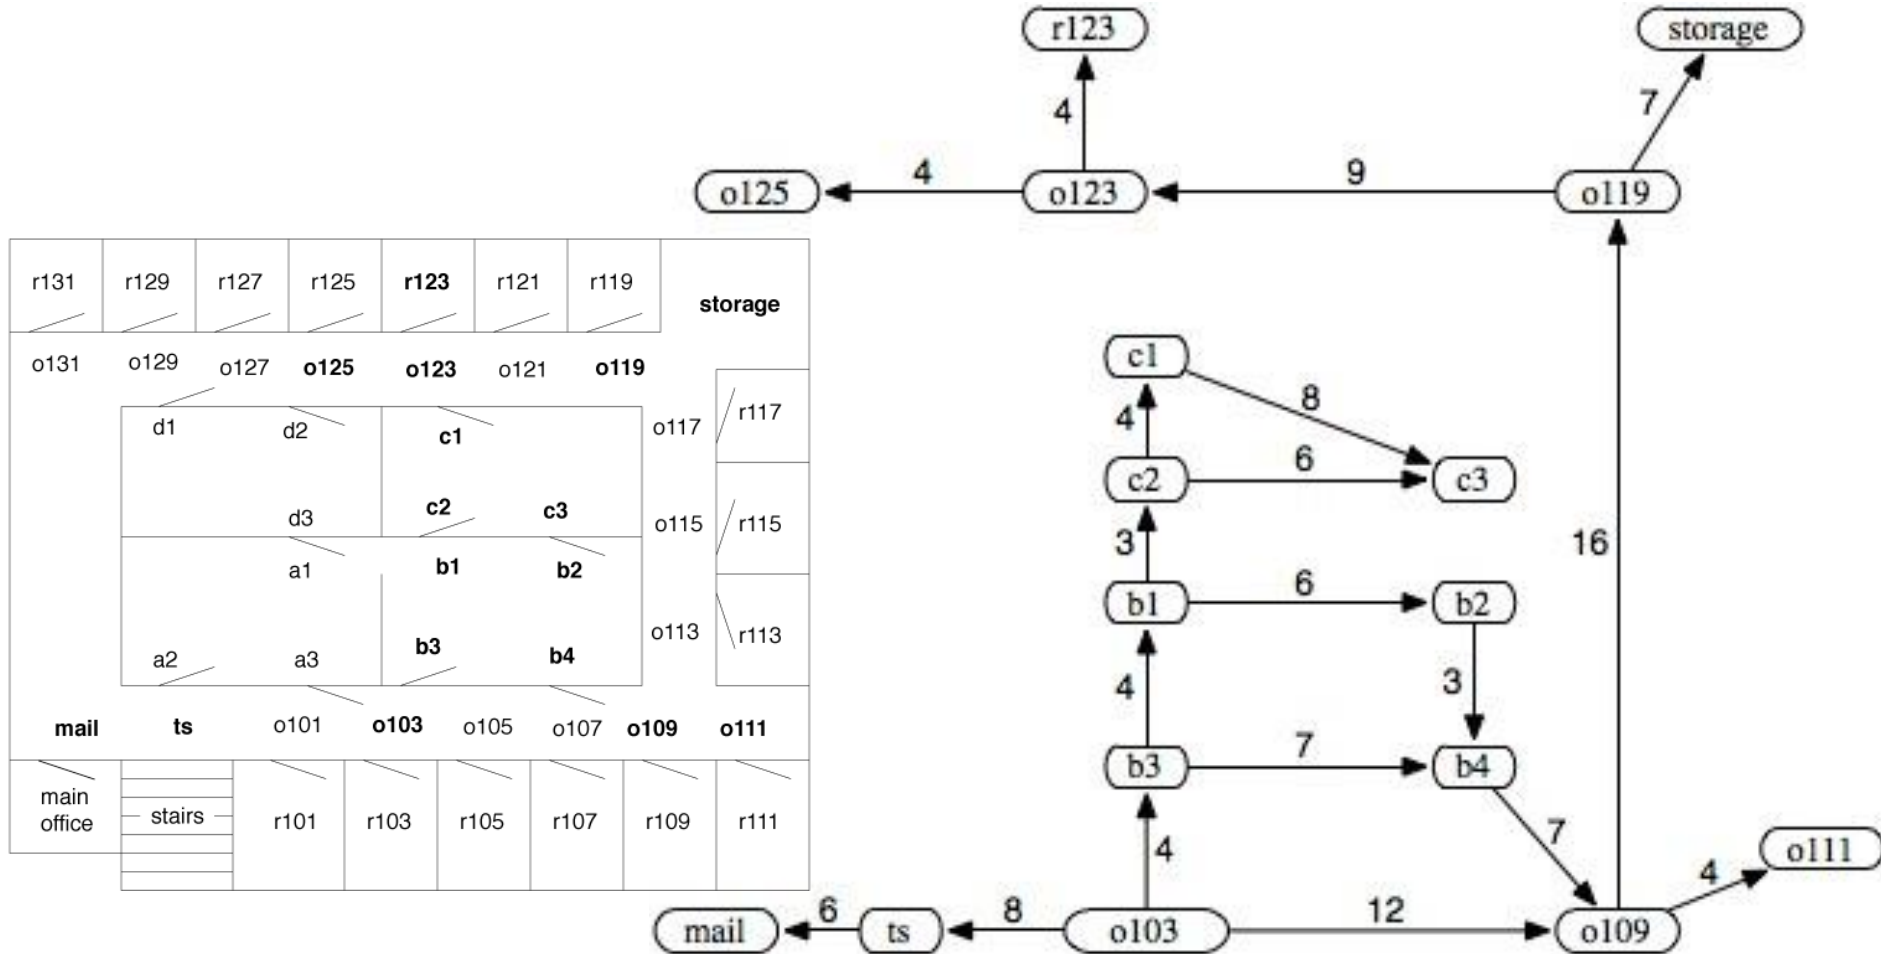
Environment for Delivery Robot (ch.1.6)

Simplified

- Consider only bold locations here
- Limits in direction of movement (e.g., can only move in the direction a door opens, can't go back to previous location, etc.)
- Start: o103
- Goal: r123



Search Space for the Delivery Robot



Learning Goals for today's class

- **Define** what is a representation and reasoning system
- **Differentiate** between single/static and sequential problems, as well as between deterministic and stochastic ones

TO DO for next class

- Read Ch 3 (3.1-3.5.2, 3.7.3)
- Assignment 0

• Review the definitions in the next three slides. I won't go over them next week

Graphs

- A **directed graph** consists of a set N of **nodes (vertices)** and a set A of ordered pairs of nodes, called **edges (arcs)**.
- Node n_2 is a **neighbor** of n_1 if there is an arc from n_1 to n_2 . That is, if $\exists n_1, n_2 \in A$.

- A **path** is a sequence of nodes n_0, n_1, \dots, n_k such that $n_{i-1}, n_i \in A$.
- A **cycle** is a non-empty path such that the start node is the same as the end node.
- A **directed acyclic graph** (DAG) is a graph with no cycles
- Given a set of start nodes and goal nodes, a **solution** is a path from a start node to a goal node


Graph specification for the Delivery Robot

$N = \{\text{mail}, \text{ts}, \text{o103}, \text{b3}, \text{o109}, \dots\}$

$A = \{$
 $\langle \text{ts}, \text{mail} \rangle,$
 $\langle \text{o103}, \text{ts} \rangle, \langle \text{o103}, \text{b3} \rangle,$
 $\langle \text{o103}, \text{o109} \rangle,$
 $\dots\}$

One of several solution paths:

$\langle \text{o103}, \text{o109}, \text{o119}, \text{o123}, \text{r123} \rangle$



Branching Factor

- The **forward branching factor** of a node is the number of arcs going out of the node
- The **backward branching factor** of a node is the number of arcs going into the node
- If the forward branching factor of a node is **b** and the graph is a tree, how many **nodes** are **n steps away** from that node?