

Lecture 6

Analysis of A^* ,

B&B

Search

Refinements

(Ch 3.7.1 - 3.7.4) Slide 1

Lecture Overview

- ➔ Recap of previous lecture
 - Analysis of A^*
 - Branch-and-Bound

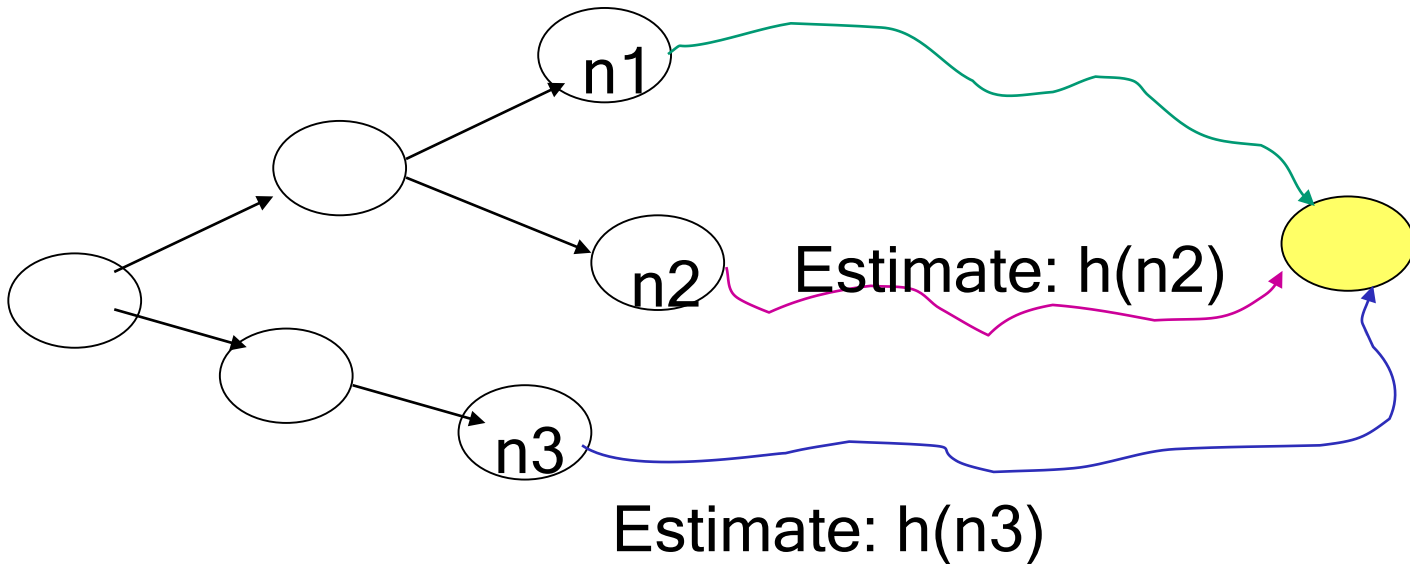
- Cycle checking, multiple path pruning
- Stored Graph - Dynamic Programming

Slide 2

How to Make Search More Informed?

Def.: A **search heuristic $h(n)$** is an estimate of the cost of the optimal (cheapest) path from node n to a goal node.

Estimate: $h(n_1)$




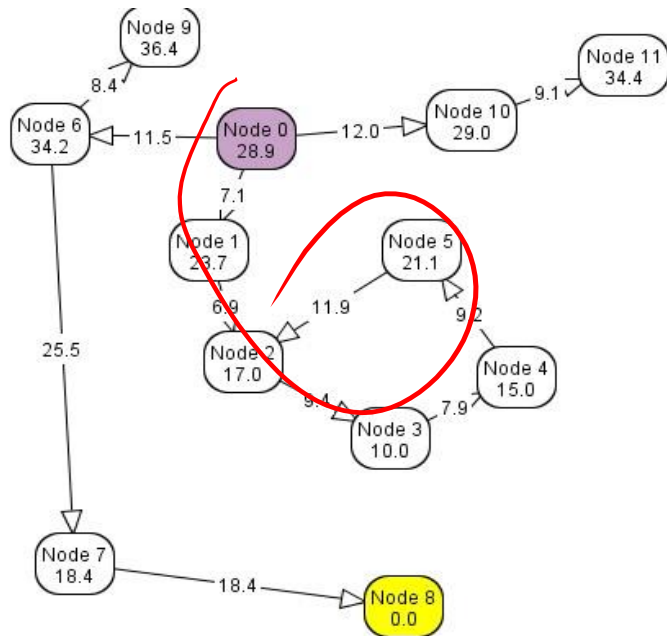
- h can be extended to paths: $h([n_0, \dots, n_k]) = h(n_k)$
- $h(n)$ should leverage readily obtainable information (easy to compute) about a node.

Best First Search (BestFS)

- Always choose the path on the frontier with the smallest h value.
- BestFS treats the frontier as a priority queue ordered by h .
- Can get to the goal pretty fast if it has a good h but...

It is not complete,

nor optimal



- still has time and space worst-case complexity of $O(b^m)$

4

Learning Goal for Search

Apply basic properties of search algorithms:

- completeness, optimality, time and space complexity

	Complete	Optimal	Time	Space
DFS	N (Y if no cycles)	N	$O(b^m)$	$O(mb)$
BFS	Y	Y	$O(b^m)$	$O(b^m)$
IDS	Y	Y	$O(b^m)$	$O(mb)$
LCFS (when arc costs available)	Y Costs $> \epsilon$	Y Costs ≥ 0	$O(b^m)$	$O(b^m)$
Best First (when available)	N	N	$O(b^m)$	$O(b^m)$

uninformed



 Uninformed but using arc cost

 Informed (goal directed)

Remind definition of admissible... Example 3: Eight Puzzle

- Another possible $h(n)$:

Sum of number of moves between each tile's current position and its goal position (we can move over other tiles in the grid)




Sum (

Example 3: Eight Puzzle

- Another possible $h(n)$:

Sum of number of moves between each tile's current position and its goal position



1 2 3 4 5 6 7 8

$$\text{sum } (2 \ 3 \ 3 \ 2 \ 4 \ 2 \ 0 \ 2) = 18$$

Admissible?


A. Yes

B. No

C. It depends

Example 3: Eight Puzzle

- Another possible $h(n)$:
Sum of number of moves between each tile's current position
and its goal position



$$\text{sum } \begin{matrix} \boxed{1} & \boxed{2} & \boxed{3} & \boxed{4} & \boxed{5} & \boxed{6} & \boxed{7} & \boxed{8} \\ (2 & 3 & 3 & 2 & 4 & 2 & 0 & 2) \end{matrix} = 18$$

Admissible? **YES!** One needs to make at least as many moves to get to the goal state when constrained by the grid structure

How can we effectively use $h(n)$

Maybe we should combine it with the cost. How?
Shall we select from the frontier the path p with:

A. ~~Lowest $\text{cost}(p) - h(p)$~~


B. Highest $\text{cost}(p) - h(p)$

C. Highest $\text{cost}(p) + h(p)$

D. ~~Lowest $\text{cost}(p) + h(p)$~~

CPSC 322, Lecture 6

Lecture Overview

- Recap of previous lecture
- Analysis of  A*
- Branch-and-Bound
- Cycle checking, multiple path pruning
- Stored Graph - Dynamic Programming

Slide 11

A* Search Algorithm

- A* is a mix of:
- lowest-cost-first and
- best-first search
- A* treats the frontier as a priority queue ordered by $f(p) = g(p) + h(p)$

- It always selects the node on the frontier with the lowest f estimated total distance.

Analysis of A*

If the heuristic is completely uninformative and the edge costs are all the same, A* is equivalent to...

A. BFS

B. LCFS

C. DFS

D. None of the Above

Analysis of A^*

Let's assume that arc costs are strictly positive.

- Time complexity is $O(b^m)$
- the heuristic could be completely uninformative and the edge costs could all be the same, meaning that A^* does the same thing as...
- Space complexity is $O(b^m)$ like *BFS*, A^* maintains a frontier which grows with the size of the tree

- Completeness: yes.

- Optimality: ??

Optimality of A^*

If A^* returns a solution, that solution is guaranteed to be optimal, as long as

When

- the branching factor is finite
- arc costs are strictly positive
- $h(n)$ is an underestimate of the length of the shortest path from n to a goal node, and is non-negative

Theorem

If A^* selects a path p as the solution, p is the shortest (i.e., lowest-cost) path.


CPSC 322, Lecture 8

Why is A^* optimal?


- A^* returns p
- Assume for contradiction that some other path p' is actually the shortest path to a goal

$$c(p) > c(p')$$

- Consider the moment when p is chosen from the frontier.
Some part of path p' will also be on the frontier; let's call this partial path p'' .



$cost(p')$



$$C(p') < C(p)$$

$$f(p) \leq f(p'')$$

$$C(p) + h(p) \leq C(p'') + h(p'')$$

$$C(p) \leq C(p'') + h(p'')$$

- Because p was expanded before p' ,
- Because p is a goal, $h(p) = 0$. Thus
- Because h is admissible, $cost(p'') + h(p'') \leq C(p')$ for any path p' to a goal that extends p''
- Thus $C(p) \leq C(p')$ for any other path p' to a goal.

Why is A optimal? (cont*)

This contradicts our assumption that p is the shortest path.

Optimal efficiency of A^*

- In fact, we can prove something even stronger about A^* : in a sense (given the particular heuristic that is available) no search algorithm could do better!
- **Optimal Efficiency:** Among all optimal algorithms that start from the same start node and use the

same heuristic, A^* expands the minimal number of paths.

Samples A^* applications

- An Efficient A^* Search Algorithm For Statistical Machine Translation. 2001
- The Generalized A^* Architecture. Journal of Artificial Intelligence Research (2007)
- Machine Vision ... Here we consider a new compositional model for finding salient curves.
- Factored A^* search for models over sequences and trees International Conference on AI. 2003....

It starts saying... The primary challenge when using A* search is to find heuristic functions that ~~simultaneously are admissible, close to actual completion costs, and efficient to calculate...~~ *(Natural Language Processing)* applied to NLP and BioInformatics

CPSC 322, Lecture 9

Slide 19

Samples A* applications (cont')

Aker, A., Cohn, T., Gaizauskas, R.: Multi-document summarization using A* search and discriminative training. Proceedings of the 2010 Conference on Empirical Methods in Natural Language

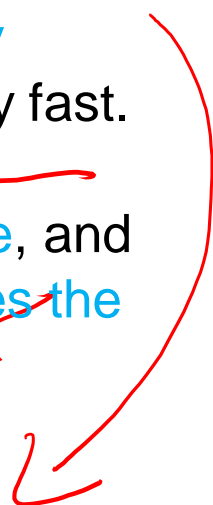
Processing.. ACL (2010)

Samples A* applications (cont')

EMNLP 2014 A* CCG Parsing with a

Supertagfactored Model M. Lewis, M. Steedman

We introduce a new CCG parsing model which is factored on lexical category assignments. Parsing is then simply a deterministic search for the most probable category sequence that supports a CCG derivation. The parser is extremely simple, with a tiny feature set, no POS tagger, and no statistical model of the derivation or dependencies. Formulating the model in this way allows a highly effective heuristic for A* parsing, which makes parsing extremely fast. Compared to the standard C&C CCG parser, our model is more accurate out-of-domain, is four times faster, has higher coverage, and is greatly simplified. We also show that using our parser improves the performance of a state-of-the-art question answering system



Follow up ACL 2017 (main NLP conference – was held in Vancouver in August!) A* CCG Parsing with a Supertag and Dependency Factored

Slide 21

Model Masashi Yoshikawa, Hiroshi CPSC 322, Lecture 8 Noji, Yuji
Matsumoto

A* advantages

What is a key advantage of A* ?

- A. Does not need to consider the cost of the paths
- B. Has a linear space complexity
- C. It is often optimal
- ☒ D. None of the above

CPSC 322, Lecture 6

Lecture Overview

➡ Recap of previous lecture • Analysis of A^* • Branch-and-Bound • Cycle checking, multiple path pruning • Stored Graph - Dynamic Programming

Slide 23

Branch-and-Bound Search

- Biggest advantages of A*...

$$\underline{f = c + h}$$

- What is the biggest problem with A*?


b^m

- Possible, preliminary ~~Solution:~~

f

Branch-and-Bound Search Algorithm

- Follow exactly the same search path as **depth-first search**
 - treat the frontier as a stack: expand the most-recently added path first
 - the order in which neighbors are expanded can be governed by some arbitrary node-ordering heuristic



we can use
 $f = c + h$

Once this strategy has found a solution...

What should it do next ?

- A. Keep running DFS, looking for deeper solutions?
- ☒ B. Stop and return that solution
- C. Keep searching, but only for shorter solutions

D. None of the above


CPSC 322, Lecture 6

Branch-and-Bound Search Algorithm

- Keep track of a lower bound and upper bound on solution cost at each path
- lower bound: $LB(p) = f(p) = cost(p) + h(p)$
- upper bound: $UB = \text{cost of the best solution found so far.}$
 - ✓ if no solution has been found yet, set the upper bound to ∞ .
- When a path p is selected for expansion:
- if $LB(p) > UB$, remove p from frontier without expanding it (pruning)

- else expand p, adding all of its neighbors to the frontier

The numbers correspond to f for the path from start to that node




CPSC 322, Lecture 9

$UB = \infty$
 \uparrow
 6
 4
 Same for all paths at any given time

Branch-and-Bound Search Algorithm

- Keep track of a **lower bound** and **upper bound** on solution cost at each path
- **lower bound**: $LB(p) = f(p) = \text{cost}(p) + h(p)$
- **upper bound**: $UB = \text{cost of the best solution found so far.}$

- ✓ if no solution has been found yet, set the upper bound to ∞ .
- When a path p is selected for expansion:
 - if $LB(p) \geq UB$, remove p from frontier without expanding it (pruning)
 - else expand p , adding all of its neighbors to the frontier



$$\underline{UB} = \frac{\infty}{64}$$

CPSC 322, Lecture 9

- Arc cost = 1

-

Before expanding a path p ,

$h(n) = 0$ for every n

JUST TO SIMPLIFY
THE EXAMPLE

1


check its f value $f(n)$:

- Arc cost = 1

-

Before expanding a path p ,

$h(n) = 0$ for every n




check its f value $f(p)$:
Expand only if $f(p) \leq UB$

- Arc cost = 1

-

Before expanding a path p ,

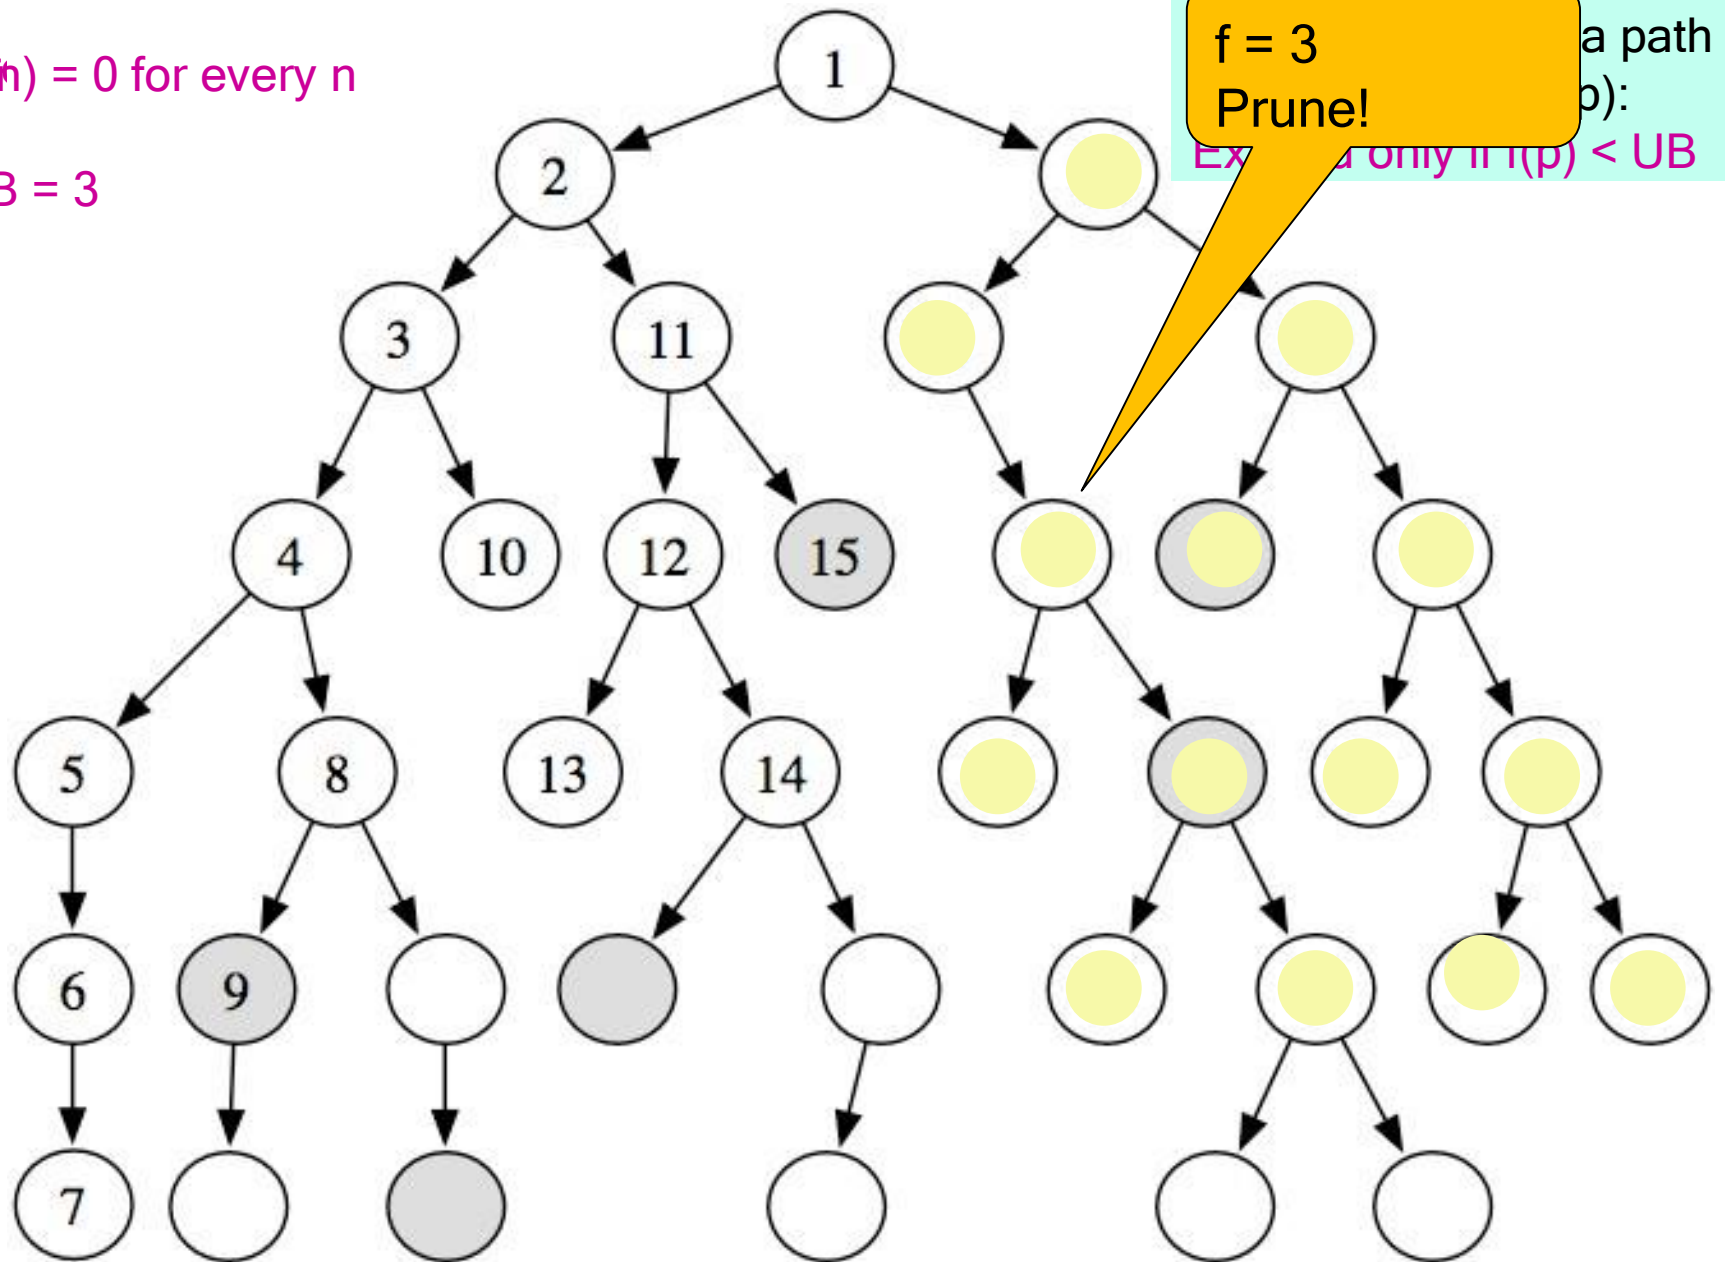
$h(n) = 0$ for every n



check its f value $f(p)$;
Expand p if $f(p) \leq \text{UB}$

Solution!

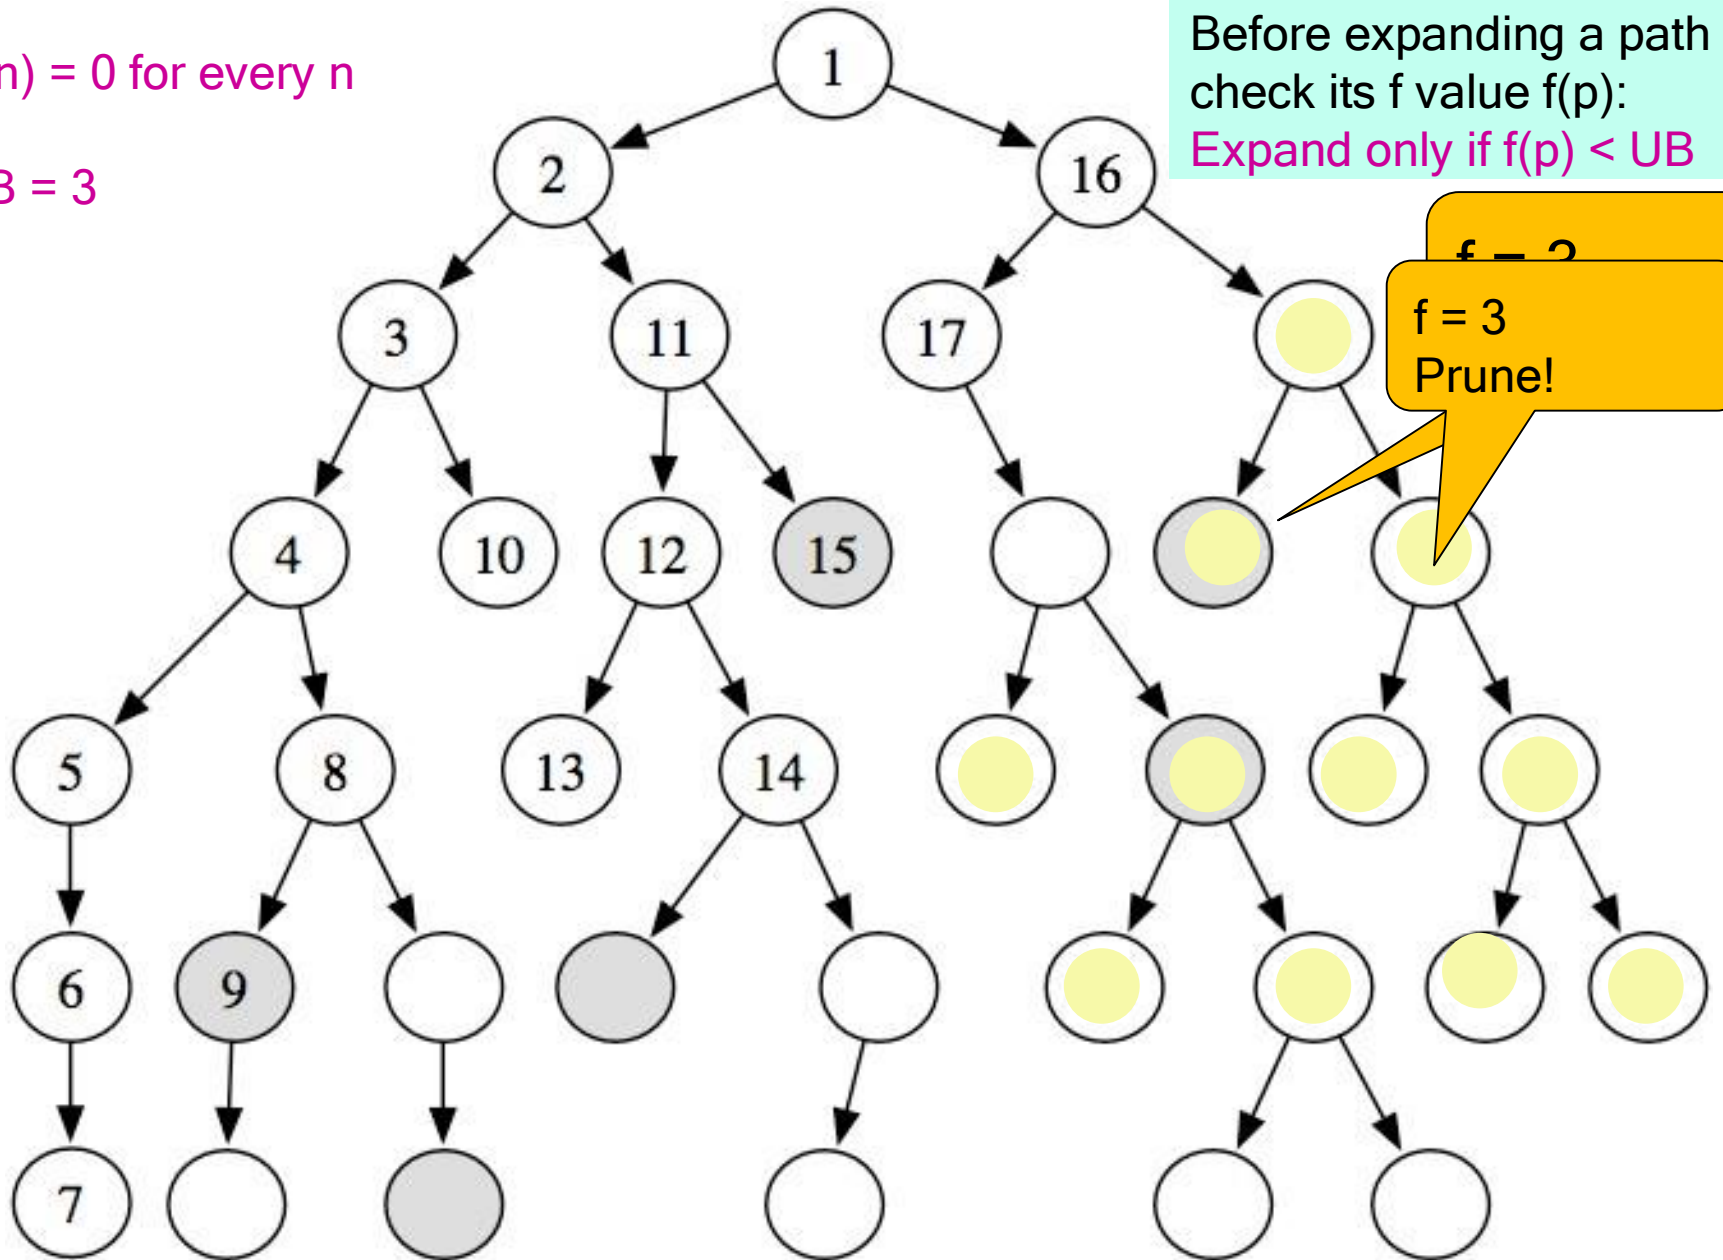
- $f(n) = 0$ for every n
-
- $UB = 3$



$h(n) = 0$ for every n

- UB = 3

Before expanding a path p ,
check its f value $f(p)$:
Expand only if $f(p) < UB$



- Arc cost = 1

-

Branch-and-Bound Analysis

- Completeness: _____
- however, for many problems of interest _____
- Time complexity: $O(b^m)$
- Space complexity: b^m
- Branch & Bound has the same space complexity as DTS.



- this is a big improvement over

A*

_____!

- Optimality: _____


CPSC 322, Lecture 9

Lecture Overview

- Recap of previous lecture
 - Analysis of A^*
 - Branch-and-Bound
 - Cycle checking, multiple path pruning
 - Stored
- Graph - Dynamic Programming

Slide 35

Cycle Checking



You can prune a path that ends in a node already on the path.
This pruning cannot remove an optimal solution.



- The time for checking is

_____ in path length.




Repeated States / Multiple Paths

Failure to detect repeated states can turn a linear problem into an exponential one!



Multiple-Path Pruning




- You can prune a path to node n that you have already found a path to
- (if the new path is longer – more costly).

Multiple-Path Pruning & Optimal Solutions

Problem: what if a subsequent path to n is shorter than the first path to n ?


- You can remove all paths from the frontier that use the longer path. (as these can't be optimal)



Multiple-Path Pruning & Optimal Solutions

Problem: what if a subsequent path to n is shorter than the first path to n ?

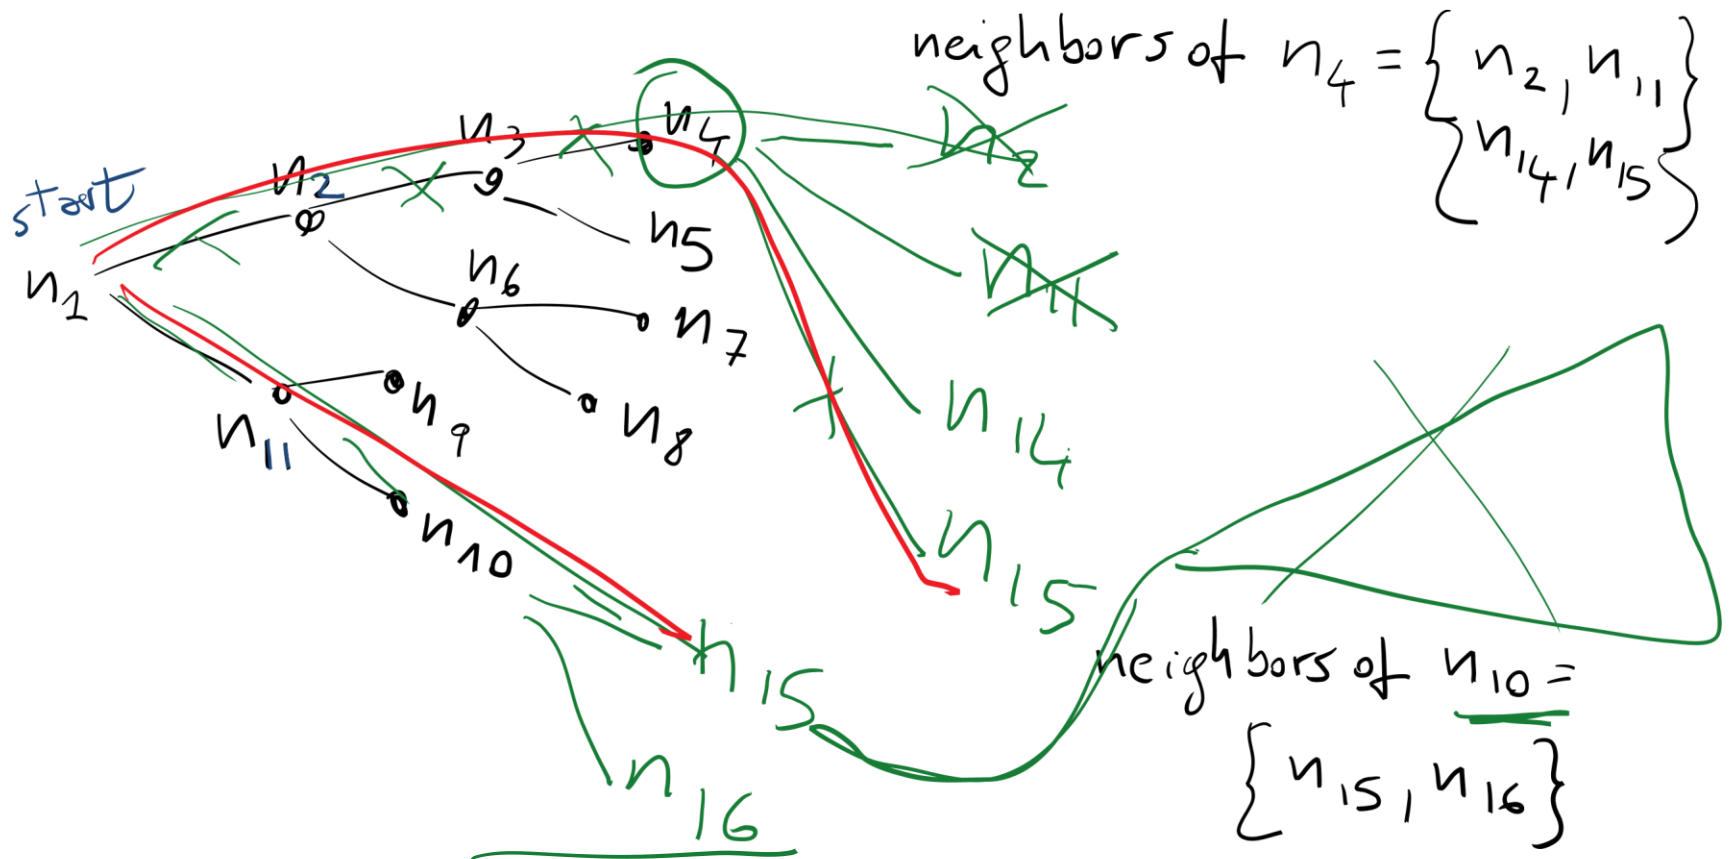
- You can change the initial segment of the paths on the frontier to use the shorter path.



Example

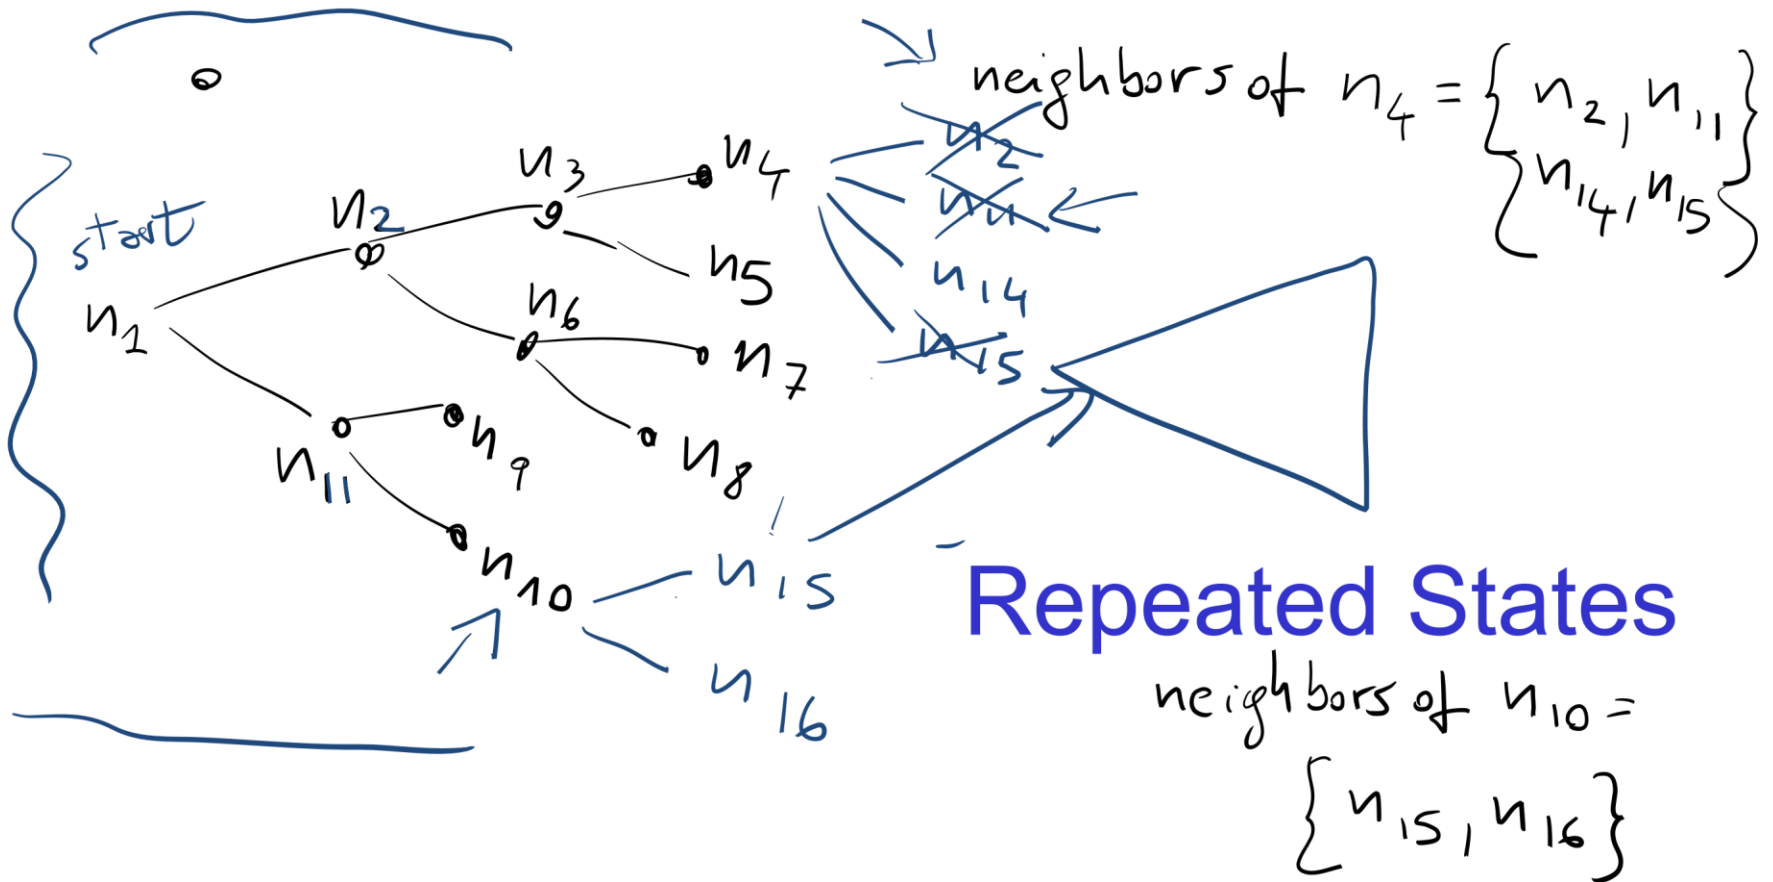
Pruning Cycles

Repeated States




Example

Pruning Cycles



Lecture Overview

Dynamic Programming


- Recap of previous lecture
- Analysis of A^*
- Branch-and-Bound
- Cycle checking, multiple path pruning 
- Stored Graph - Dynamic Programming

Dynamic Programming

- Idea: for statically stored graphs, build a table of $\text{dist}(n)$:
- The **actual distance** of the shortest path from node n to a goal g
- This is the perfect_____.

Dynamic Programming

- dist(g) = 0
- dist(z) = 1
- dist(c) = 3
- dist(b) = 4
- dist(k) = ?
- dist(h) = ?



- How could we implement that?


This can be built **backwards** from the goal:



Dynamic Programming

all the neighbors m

$$dist(n) = \begin{cases} 0 & \text{if } \text{is_goal}(n), \\ \min_{\langle n, m \rangle \in A} \left[\text{cost}(n, m) + dist(m) \right] & \text{otherwise} \end{cases}$$



dist (g)

0

dist (b)

$$\min[(2 + 0)] = 2$$

dist (c)

$$\min[(3 + 0)] = 3$$

dist (a)

A. $\min(3, 3)$

B. $\min(6, 3)$

dist(a)

iClicker.



C.min(2,3)

Dynamic Programming

This can be built **backwards** from the goal:

This can be built **backwards** from the goal:

□□ 0 if is goal(n).

$$\text{dist}(n) = \begin{cases} 0 & \text{if } n = g \\ \min_{m \in A} (\text{cost}(n, m) + \text{dist}(m)) & \text{otherwise} \end{cases}$$


all the neighbors m

$$\text{dist}(b) = \min[2 + 0] = 2$$

$$\text{dist}(c) = \min[3 + 0] = 3$$

$$\text{dist}(a) = \min[(3 + 3), (1 + 2)] = 3$$

3CPSC 322, Lecture 9




Dynamic Programming


This can be used locally to determine what to do.

From each node n go to its neighbor which minimizes

$$\square \text{cost}(n,m) \square \text{dist}(m) \square$$

But there are at least two main problems



- You need enough space to store the graph. 
- The dist function needs to be recomputed for each goal

CPSC 322, Lecture 9

Learning Goals for today's class

- Define/read/write/trace/debug & Compare different Informed search algorithms Best First Search, A^* , and Branch Bound
- Formally prove A^* optimality.

Dynamic Programming

- Apply techniques to deal with cycles and repeated states
- Simplify search when full search graph can be stored

CPSC 322, Lecture 7

To Do for Next Class

- Read
 - Chp 4.1-4.2 (Intro to Constraint Satisfaction Problems)
-
- Do Practice Exercise 3E
 - Keep working on assignment-1 !

Next class

Finish Search (finish Chpt 3)

- Branch-and-Bound
- Informed IDS
- A* enhancements
- Non-heuristic Pruning
- Dynamic Programming

